

Logistic Regression

Ishita Jaju
16BCE1059

1 The algorithm

Logistic Regression is based on the Gradient Descent algorithm. Gradient descent is an algorithm to minimize the cost function of a particular dataset and find the required weights of the hypothesis.

Gradient descent works on simultaneously updating the value for every weight like so:

Repeat until convergence {

$$\theta_j \leftarrow \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta)$$

}

where θ are the weights of the hypothesis, J is the cost function to minimize and α is the learning rate.

For logistic regression, the hypothesis function is the sigmoid function, with the input being $g(\theta^T x)$ (where g is the sigmoid function). So the hypothesis would be:

$$h_{\theta}(x) = \frac{1}{1 + e^{-\theta^T x}}$$

where θ is the weight vector and x is the attribute matrix.

1.1 Logistic Regression for multiclass classification

The above equation of the hypothesis gives an output between 0 and 1. So classes are defined based on the nearest integer between 0 and 1, which is binary classification. For multiclass classification, the logistic regression classifier is trained as $h_{\theta}^{(i)}(x)$ for each class i to predict the probability that $y = i$. So, on a new input x , we need to make a prediction and pick the class i that maximises the hypothesis: $\max_i h_{\theta}^{(i)}(x)$

2 Using the seeds dataset

2.1 The dataset


```
headers = ["area","perim","compactness","kernel_length","kernel_width","ass_coeff","kernel_grove_length","class"]
seeds = pd.read_csv("/home/ijaju/Documents/datasets/seeds_dataset.csv",header = None,names=headers)
seeds.tail()
```

	area	perim	compactness	kernel_length	kernel_width	ass_coeff	kernel_grove_length	class
205	12.19	13.20	0.8783	5.137	2.981	3.631	4.870	3
206	11.23	12.88	0.8511	5.140	2.795	4.325	5.003	3
207	13.20	13.66	0.8883	5.236	3.232	8.315	5.056	3
208	11.84	13.21	0.8521	5.175	2.836	3.598	5.044	3
209	12.30	13.34	0.8684	5.243	2.974	5.637	5.063	3

The dataset has attributes like area,kernel grove length etc and the classification is done for 3 species of wheat. The species have 70 entries each and there are no null values in any of the columns.

```
sb.countplot(x='class',data=seeds, palette='hls')
```

<matplotlib.axes._subplots.AxesSubplot at 0x7fefaf9f97b8>



```
seeds.isnull().sum()
```

```
area          0
perim         0
compactness   0
kernel_length 0
kernel_width  0
ass_coeff     0
kernel_grove_length 0
class         0
dtype: int64
```

2.2 Running the algorithm

On running the algorithm, the accuracy comes out to be 91% for this multiclass dataset.

```
LogReg = LogisticRegression()
LogReg.fit(X_train, y_train)

y_pred = LogReg.predict(X_test)

confusion_matrix = confusion_matrix(y_test, y_pred)
confusion_matrix
array([[17,  3,  2],
       [ 1, 25,  0],
       [ 0,  0, 15]])

print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
1	0.94	0.77	0.85	22
2	0.89	0.96	0.93	26
3	0.88	1.00	0.94	15
avg / total	0.91	0.90	0.90	63

3 My dataset - Chronic Kidney Disease

This is a binary classification dataset with the classes being - 'ckd' and 'notckd'.

```
ckd.dtypes
age          float64
bp           float64
sg           float64
al           float64
su           float64
rbc          object
pc           object
pcc          object
ba           object
bgr          float64
bu           float64
sc           float64
sod          float64
pot          float64
hemo         float64
pcv          int64
wc           int64
rc           float64
htn          object
dm           object
cad          object
appet        object
pe           object
ane          object
classification object
dtype: object
```

With there being so many 'object' datatypes where the inputs are classifiers, like 'normal' and 'abnormal'; and 'present' and 'notpresent', we need to convert these into numerical inputs like '0' and '1' or '-1' and '1' for the algorithm to work on them. The function below takes care of the same:

```
class_types = ['rbc', 'pc', 'pcc', 'ba', 'htn', 'dm', 'cad', 'appet', 'pe', 'ane']
for i in ckd.columns:
    for j in class_types:
        if i==j:
            e = ckd[j].unique()
            ckd[j] = np.where(ckd[j]==e[0],-1,1)
```

Then after running the algorithm, the accuracy is found to be 97%.

```
X = ckd.iloc[:,0:23].values
y = ckd.iloc[:,24].values
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = .3, random_state=25)
LogReg = LogisticRegression()
LogReg.fit(X_train, y_train)
```

```
LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
                    intercept_scaling=1, max_iter=100, multi_class='ovr', n_jobs=1,
                    penalty='l2', random_state=None, solver='liblinear', tol=0.0001,
                    verbose=0, warm_start=False)
```

```
y_pred = LogReg.predict(X_test)
```

```
from sklearn.metrics import accuracy_score
accuracy_score(y_test, y_pred)
```

```
0.9791666666666666
```